



EDB

Postgres® for the AI Generation

EXPLAIN Explained: Making Sense of PostgreSQL Query Plans

April 25th, 2025. PGDay Pune

Jeevan Chalke

Principal Engineer, Engineering Manager

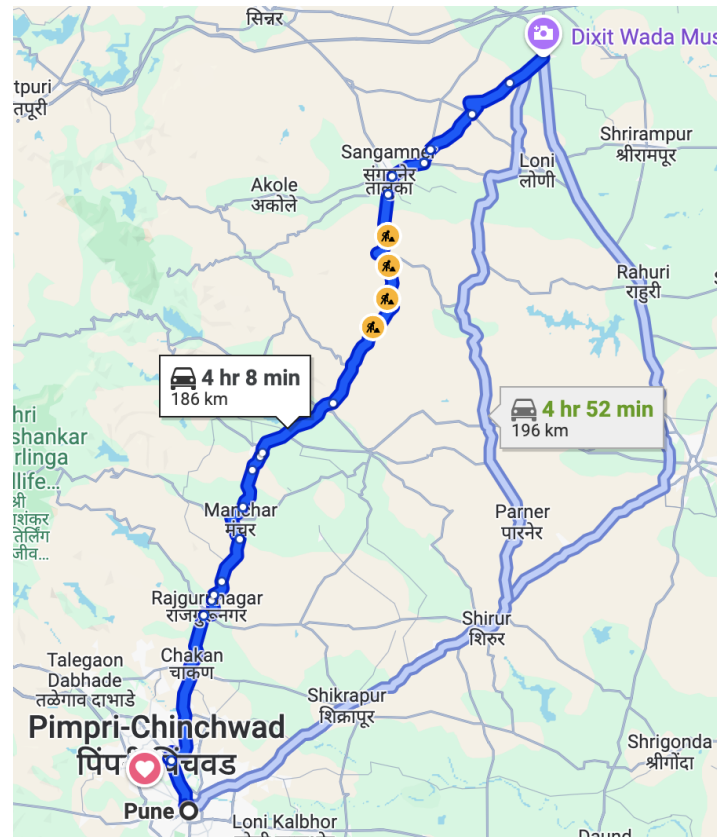
Jeevan.chalke@enterprisedb.com

Agenda

- Brief on optimizer/planner
- Paths and Plan
- EXPLAIN and EXPLAIN ANALYZE
 - Scans - seq, index
 - Joins
 - Sort
 - Grouping
 - Parallelism, Partitioning, Foreign Scan - Summary
- Options used with EXPLAIN

Brief on optimizer/planner

- Statistics
 - ANALYZE table
- Creates different Paths
 - The way query can be evaluated to give the desired output
- Gives metrics to them
- Chooses the cheapest path
- Converts Path to the Plan
- Executor then executes the query according to the chosen plan



Paths and Plan

- Costs
 - cpu_tuple_cost
 - cpu_operator_cost
 - ...
- Rows
 - Estimated
 - Actual
- Width
- Time
 - Actual
 - Planning
 - Execution

EXPLAIN works with SELECT, INSERT, UPDATE, DELETE, VALUES, EXECUTE, DECLARE, CTAS, or CREATE MATERIALIZED VIEW queries

Basic EXPLAIN with seq scan, setup

```
# create table mytab1(a int, b int, c int);
# create table mytab2(x int, y int);
# insert into mytab1 select i, i%50, i%100 from generate_series(1, 1000000) i;
# insert into mytab2 select i, i%30 from generate_series(1, 1000) i;
# analyze mytab1;
# analyze mytab2;
```

```
# select relname, relpages, reltuples from pg_class where relname like 'mytab%';
 relname | relpages | reltuples
-----+-----+-----
 mytab1  |      5406 |      1e+06
 mytab2  |         5 |       1000
(2 rows)
```

```
# select name, setting from pg_settings where name in ('cpu_tuple_cost', 'seq_page_cost');
      name      | setting
-----+-----
cpu_tuple_cost  | 0.01
seq_page_cost   | 1
(2 rows)
```

Basic EXPLAIN with seq scan

- Syntax
 - EXPLAIN [(option [, ...])] statement
 - EXPLAIN [ANALYZE] [VERBOSE] statement

```
# explain
select * from mytab1;
```

QUERY PLAN

Seq Scan on mytab1 (cost=0.00..**15406.00** rows=**1000000** width=12)
(1 row)

- For a seq scan; the planner has to do two things
 - Read all the pages
 - Read all the tuples from each page
- So the cost will be
 - $5406 \text{ (relpages)} * 1 \text{ (seq_page_cost)} + 1000000 \text{ (reltuples)} * 0.01 \text{ (cpu_tuple_cost)} =$
15406.00

With WHERE clause and INDEX

```
# explain select * from mytab1 where a < 800000;
```

QUERY PLAN

```
Seq Scan on mytab1 (cost=0.00..17906.00 rows=798548 width=12)
```

```
  Filter: (a < 800000)
```

```
(2 rows)
```

- Filter is added in the output showing the condition; Number of rows reduced
- Cost increased: `cpu_operator_cost` = 0.0025 (+2500)
- Let's add an Index and run the same query

```
# create index myidx1 on mytab1(a);
```

```
# explain select * from mytab1 where a < 800000;
```

QUERY PLAN

```
Seq Scan on mytab1 (cost=0.00..17906.00 rows=798548 width=12)
```

```
  Filter: (a < 800000)
```

```
(2 rows)
```

Same result?

Let's understand
EXPLAIN ANALYZE
first...

EXPLAIN ANALYZE

- Gives the EXPLAIN plan
- Executes the query; discards the output
- Gives actual timing and row details
- Gives a few more finer details; like loops and allows using BUFFERS option
- Provides summary showing planning and execution time
- Risky with DML commands
 - For example, an EXPLAIN ANALYZE DELETE... will actually delete the rows from a table

With an Index and EXPLAIN ANALYZE

```
# explain (analyze) select * from mytab1 where a < 800000;  
QUERY PLAN
```

```
Seq Scan on mytab1 (cost=0.00..17906.00 rows=799624 width=12) (actual  
time=0.019..134.772 rows=799999 loops=1)  
  Filter: (a < 800000)  
    Rows Removed by Filter: 200001  
  Planning Time: 0.085 ms  
  Execution Time: 170.264 ms  
(5 rows)
```

```
# set enable_seqscan to off; -- Forces Index scan  
# explain (analyze) select * from mytab1 where a < 800000;  
QUERY PLAN
```

```
Index Scan using myidx1 on mytab1 (cost=0.42..27099.84 rows=799624 width=12) (actual  
time=0.045..251.022 rows=799999 loops=1)  
  Index Cond: (a < 800000)  
  Planning Time: 0.083 ms  
  Execution Time: 285.570 ms  
(4 rows)
```

With an Index and EXPLAIN ANALYZE (continued...)

```
# set enable_seqscan to on;
```

```
# explain (analyze) select * from mytab1 where a > 800000;  
QUERY PLAN
```

```
-----  
Index Scan using myidx1 on mytab1 (cost=0.42..6797.99 rows=200375 width=12) (actual  
time=0.070..62.023 rows=200000 loops=1)  
  Index Cond: (a > 800000)  
    Planning Time: 0.177 ms  
    Execution Time: 69.694 ms  
(4 rows)
```

- So, having an index doesn't always improve the performance.
- But it does in most of the cases and you have to figure it out by looking at your data.

With an Index, and specific column

```
# explain (analyze) select a from mytab1 where a > 800000;  
QUERY PLAN
```

```
-----  
Index Only Scan using myidx1 on mytab1 (cost=0.42..5710.99 rows=200375 width=4) (actual  
time=0.051..31.107 rows=200000 loops=1)  
  Index Cond: (a > 800000)  
  Heap Fetches: 75  
  Planning Time: 0.091 ms  
  Execution Time: 38.271 ms  
(5 rows)
```

- Index Scan is changed to Index Only Scan
- Required column is part of an index itself so no need to scan the actual table
- Available planner options
 - enable_seqscan, enable_indexscan, enable_indexonlyscan, enable_bitmapscan

JOINS

```
# explain (analyze, costs off) select a,x from mytab1 t1 left join mytab2 t2 on (t1.a=t2.x);  
QUERY PLAN
```

```
-----  
Hash Left Join (actual time=0.316..281.785 rows=1000000 loops=1)  
  Hash Cond: (t1.a = t2.x)  
    -> Seq Scan on mytab1 t1 (actual time=0.014..88.353 rows=1000000 loops=1)  
    -> Hash (actual time=0.288..0.290 rows=1000 loops=1)  
          Buckets: 1024  Batches: 1  Memory Usage: 44kB  
          -> Seq Scan on mytab2 t2 (actual time=0.013..0.132 rows=1000 loops=1)  
Planning Time: 0.572 ms  
Execution Time: 321.697 ms
```

```
# explain (costs off)select a from mytab1 t1 where a in (select x from mytab2);  
QUERY PLAN
```

```
-----  
Merge Semi Join  
  Merge Cond: (t1.a = mytab2.x)  
    -> Index Only Scan using myidx1 on mytab1 t1  
    -> Sort  
          Sort Key: mytab2.x  
          -> Seq Scan on mytab2
```

JOINS Summary

- Merge Join
 - Sorts, and then merges
 - Faster for bigger data-set
- Hash Join
 - Works with equality constraints
 - Faster provided we have enough memory
 - Mostly used where one table is smaller
- Nested Loop
 - For smaller data-set
 - Mostly used for cross joins
- Join Types: Left/Right/Full/Anti/Semi/Inner
- Planner parameters
 - `enable_hashjoin`, `enable_mergejoin`, `enable_nestloop`

SORT

```
# explain (analyze, buffers) select * from mytab1 t1 where a < 200000 order by b;  
QUERY PLAN
```

```
Sort (cost=28089.70..28594.44 rows=201894 width=12) (actual time=109.382..143.465  
rows=199999 loops=1)  
  Sort Key: b  
  Sort Method: external merge  Disk: 4320kB  
  Buffers: shared hit=1631, temp read=540 written=542  
    -> Index Scan using myidx1 on mytab1 t1 (cost=0.42..6848.57 rows=201894 width=12)  
      (actual time=0.046..61.180 rows=199999 loops=1)  
        Index Cond: (a < 200000)  
        Buffers: shared hit=1631  
Planning Time: 0.101 ms  
Execution Time: 156.481 ms  
(9 rows)
```

- Sort Key
- Sort Method is external merge sort
- Buffers option, temp read/written in blocks (8K)

SORT (Continued...)

```
# set work_mem to '16MB';
# explain (analyze, buffers, settings) select * from mytab1 t1 where a < 200000 order by b;
                                QUERY PLAN
-----
 Sort  (cost=24638.70..25143.44 rows=201894 width=12) (actual time=111.089..137.246
rows=199999 loops=1)
   Sort Key: b
   Sort Method: quicksort  Memory: 12394kB
   Buffers: shared read=1631
   ->  Index Scan using myidx1 on mytab1 t1  (cost=0.42..6848.57 rows=201894 width=12)
(actual time=3.728..82.659 rows=199999 loops=1)
     Index Cond: (a < 200000)
     Buffers: shared read=1631
Settings: enable_mergejoin = 'off', work_mem = '16MB'
Planning Time: 0.118 ms
Execution Time: 152.970 ms
(10 rows)
```

- Sort Method is quicksort due to increased **work_mem** sorting now done in-memory.

GROUPING

```
# explain (analyze, memory)
  select count(*) from mytab2 t1 group by y having sum(x) > 17000;
```

QUERY PLAN

```
HashAggregate (cost=22.50..22.88 rows=10 width=12) (actual time=0.389..0.394 rows=5
loops=1)
```

```
  Group Key: y
```

```
  Filter: (sum(x) > 17000)
```

```
  Batches: 1  Memory Usage: 24kB
```

```
  Rows Removed by Filter: 25
```

```
    -> Seq Scan on mytab2 t1 (cost=0.00..15.00 rows=1000 width=8) (actual
time=0.015..0.110 rows=1000 loops=1)
```

```
  Planning:
```

```
    Memory: used=14kB  allocated=16kB
```

```
  Planning Time: 0.132 ms
```

```
  Execution Time: 0.435 ms
```

- Group keys are displayed, and Filter shows the Having clause
- Parameter `enable_hashagg` can be used to disable HashAggregate

GROUPING with Parallelism

```
# explain verbose select count(*) from mytab3 t1 group by y having sum(x) > 17000;  
QUERY PLAN
```

```
Finalize GroupAggregate (cost=132749.06..132756.89 rows=10 width=12)
```

```
Output: count(*), y
```

```
Group Key: t1.y
```

```
Filter: (sum(t1.x) > 17000)
```

```
-> Gather Merge (cost=132749.06..132756.06 rows=60 width=20)
```

```
Output: y, (PARTIAL count(*)), (PARTIAL sum(x))
```

```
Workers Planned: 2
```

```
-> Sort (cost=131749.04..131749.11 rows=30 width=20)
```

```
Output: y, (PARTIAL count(*)), (PARTIAL sum(x))
```

```
Sort Key: t1.y
```

```
-> Partial HashAggregate (cost=131748.00..131748.30 rows=30 width=20)
```

```
Output: y, PARTIAL count(*), PARTIAL sum(x)
```

```
Group Key: t1.y
```

```
-> Parallel Seq Scan on public.mytab3 t1 (cost=0.00..94248.00  
rows=5000000 width=8)
```

```
Output: x, y
```

Parallelism, Partitioning, Foreign Scan - Summary

- Parallelism
 - Look for parameters related to Parallelism
 - `parallel_leader_participation`, `parallel_setup_cost`, `parallel_tuple_cost`, `max_parallel_workers`, `max_parallel_workers_per_gather`, etc.
 - New nodes and details in EXPLAIN
 - Gather or Gather Merge, Workers Planned
- Partitioning
 - Look for parameters related to Partitioning
 - `enable_partition_pruning`, `enable_partitionwise_aggregate`, `enable_partitionwise_join`
 - New node Append
- Foreign Scan
 - Details in EXPLAIN
 - Relations, Remote SQL

Options Supported with EXPLAIN

- **ANALYZE, VERBOSE, COSTS, SETTINGS, BUFFERS, MEMORY**
- **GENERIC_PLAN** [boolean] => Allow the statement to contain parameter placeholders like \$1, but still generate a generic plan. Cannot be used together with ANALYZE
- **SERIALIZE** [{ NONE | TEXT | BINARY }] => Include information on the cost of serializing the query's output data. Used only with ANALYZE, Default TEXT
- **WAL** [boolean] => Include information on WAL record generation. Used only with ANALYZE
- **TIMING** [boolean] => Show actual times. Used only with ANALYZE, default TRUE
- **SUMMARY** [boolean] => Gives summary on planning and execution times
- **FORMAT** { TEXT | XML | JSON | YAML } => Output display format. Default TEXT

Format

```
# explain (format json, verbose off,  
analyze, summary, timing off, costs off,  
buffers off, wal)  
select * from mytab1;  
QUERY PLAN
```

```
-----  
[                                     +  
  {                                   +  
    "Plan": {                         +  
      "Node Type": "Seq Scan",        +  
      "Parallel Aware": false,        +  
      "Async Capable": false,        +  
      "Relation Name": "mytab1",      +  
      "Alias": "mytab1",              +  
      "Actual Rows": 1000001,         +  
      "Actual Loops": 1,              +  
      "WAL Records": 0,               +  
      "WAL FPI": 0,                  +  
      "WAL Bytes": 0                  +  
    },                                +  
    "Planning Time": 0.056,           +  
    "Triggers": [                     +  
  ],                                  +  
  "Execution Time": 87.329            +  
}                                     +
```

```
# explain (format xml, verbose off,  
analyze, summary, timing off, costs off,  
buffers off, wal)  
select * from mytab1;  
QUERY PLAN
```

```
-----  
<explain xmlns="http://www.postgresql.org/2009/explain">+  
  <Query>                               +  
    <Plan>                               +  
      <Node-Type>Seq Scan</Node-Type>    +  
      <Parallel-Aware>false</Parallel-Aware> +  
      <Async-Capable>false</Async-Capable> +  
      <Relation-Name>mytab1</Relation-Name> +  
      <Alias>mytab1</Alias>              +  
      <Actual-Rows>1000001</Actual-Rows>  +  
      <Actual-Loops>1</Actual-Loops>      +  
      <WAL-Records>0</WAL-Records>        +  
      <WAL-FPI>0</WAL-FPI>               +  
      <WAL-Bytes>0</WAL-Bytes>           +  
    </Plan>                              +  
    <Planning-Time>0.105</Planning-Time>  +  
    <Triggers>                           +  
  </Triggers>                           +  
    <Execution-Time>87.991</Execution-Time> +  
  </Query>                               +  
</explain>                             +
```



EDB

Postgres® for the AI Generation

THANK YOU

References

PostgreSQL Documentation

- <https://www.postgresql.org/docs/17/sql-explain.html>
- <https://www.postgresql.org/docs/current/using-explain.html>

My same talk (PGConf India 2019)

- <https://pgconf.in/files/presentations/2019/02-0102-ExplainByJeevanChalkeEDB.pdf>

Query Planning Gone Wrong by Robert Haas

- <http://rhaas.blogspot.com/2013/05/query-planning-gone-wrong.html>

Other web links

- <https://momjian.us/main/writings/pgsql/optimizer.pdf>
- <https://neon.tech/postgresql/postgresql-tutorial/postgresql-explain>
- https://wiki.postgresql.org/images/4/45/Explaining_EXPLAIN.pdf